

CS 186 Discussion 7

Transactions, Concurrency Control
Locking Granularity

Logistics

- Nothing new...
 - Have you checked Homework 3 pushes?
 - Homework 4 due this Friday
 - Midterm 1 regrade requests due this Friday
 - Midterm 2 Thursday after break
 - Review session (probably?) Tu/Wed week of
 - Fill out Mid-Semester Feedback

Transactions

- **A**tomicity: All actions in the xact happen, or none
 - Logging
- **C**onsistency: xact will not break DB consistency
 - Integrity constraints
- **I**solation: Execution of the xact is isolated from other xacts
 - Serial ordering
- **D**urability: Committed xacts have persistent effects
 - Logging

Transactions

- **A**tomicity: All actions in the xact happen, or none
 - Logging
- **C**onsistency: xact will not break DB consistency
 - Integrity constraints
- **I**solation: Execution of the xact is isolated from other xacts
 - Serial ordering
- **D**urability: Committed xacts have persistent effects
 - Logging

Transactions

- **A**tomicity: All actions in the xact happen, or none
 - Logging
- **C**onsistency: xact will not break DB consistency
 - Integrity constraints
- **I**solation: Execution of the xact is isolated from other xacts
 - Serial ordering
- **D**urability: Committed xacts have persistent effects
 - Logging

Conflict Serializability

- What is a conflict?
 - Different xacts on one object that is *written to*

T1: R(A) R(B) W(A)

T2: R(B) W(B)

Conflict Serializability

- What is a conflict?
 - Different xacts on one object that is *written to*


T1: R(A) R(B) W(A)

T2: R(B) W(B)

Conflict Serializability

T1: R(A) R(B) W(A)


T2: R(B) W(B)



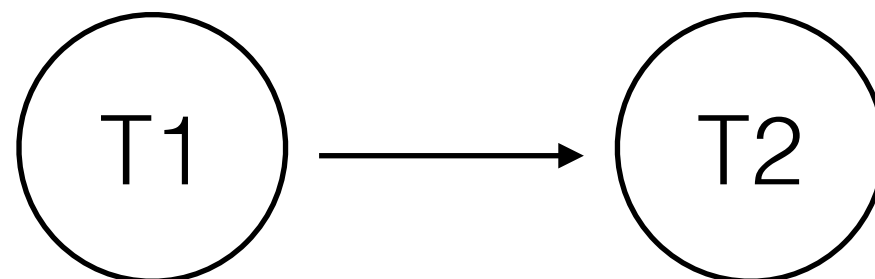
- Dependency Graph (aka Precedence Graph):
 - One node per xact
 - Edge from T_i to T_j if some operation O_i is earlier than and conflicts with O_j

Conflict Serializability

T1: R(A) R(B) W(A)
T2: R(B) W(B)



- Dependency Graph (aka Precedence Graph):
 - One node per xact
 - Edge from T_i to T_j if some operation O_i is earlier than and conflicts with O_j



Conflict Serializability

- Conflict Serializable: conflict equivalent to some serial schedule
 - Conflict Equivalent: same xacts/actions and *every conflict pair is ordered the same way*
- Serializable vs. Conflict Serializable?

Conflict Serializability

- Conflict Serializable: conflict equivalent to some serial schedule
 - Conflict Equivalent: same xacts/actions and *every conflict pair is ordered the same way*
- Serializable vs. Conflict Serializable?
 - What about view serializable?
- **Theorem**: Schedule is conflict serializable *if and only if* dependency graph is acyclic

Worksheet - 2a

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

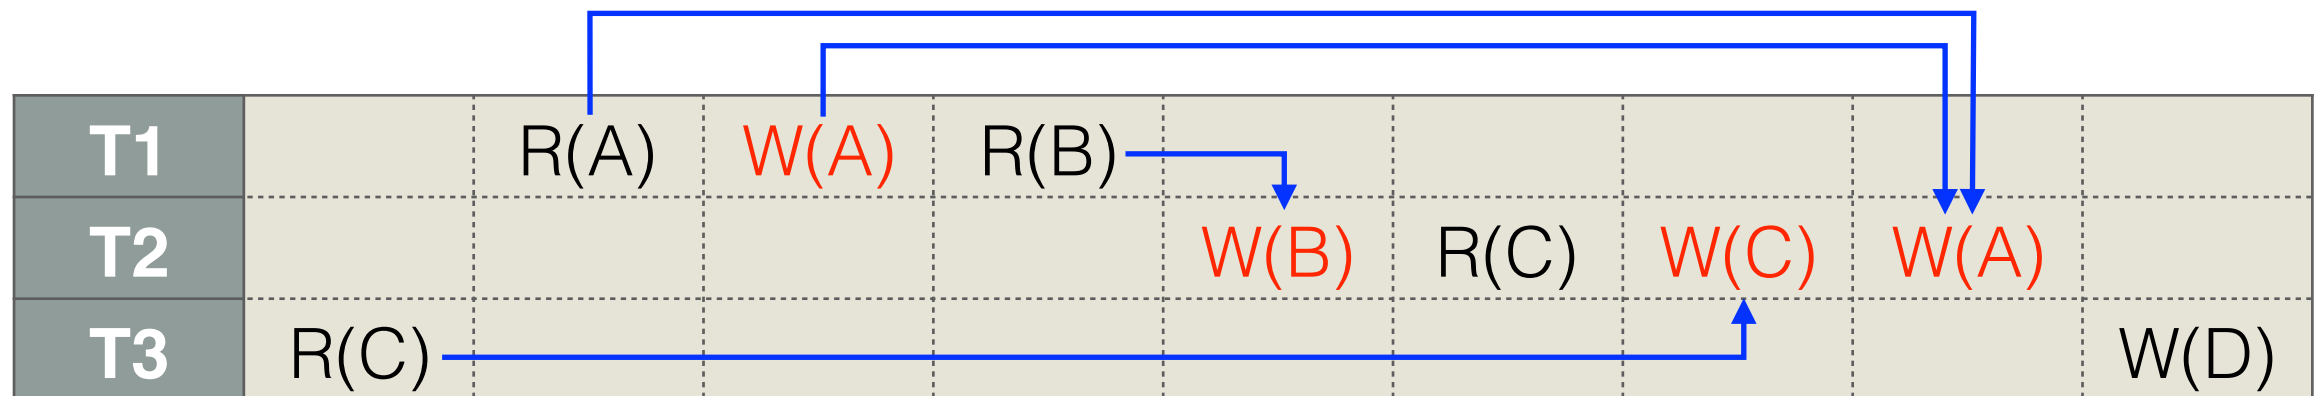
- Dependency Graph
 - Conflicting writes? Edges?

Worksheet - 2a

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

- Dependency Graph
 - Conflicting writes? Edges?

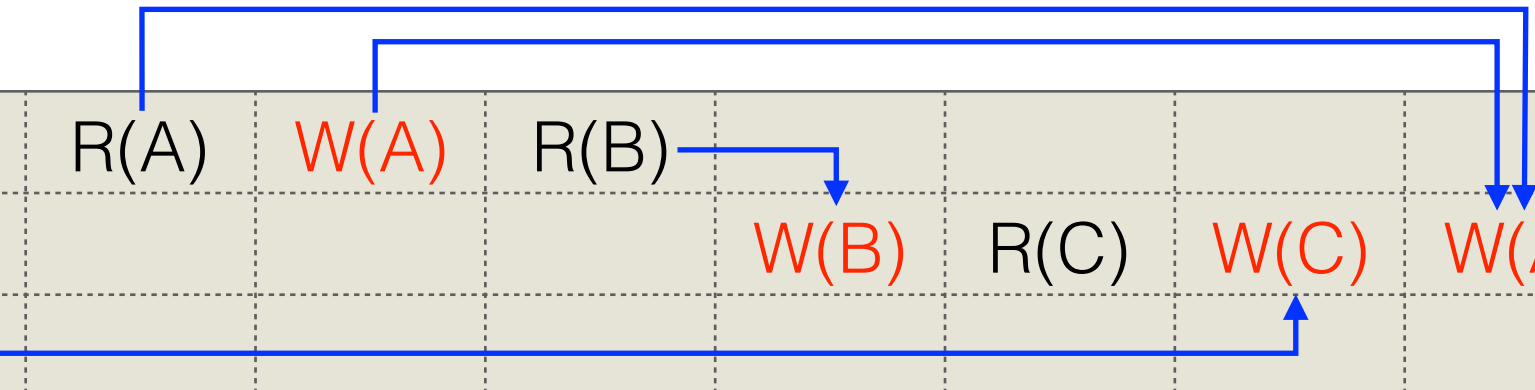
Worksheet - 2a



- Dependency Graph
 - Conflicting writes? Edges?

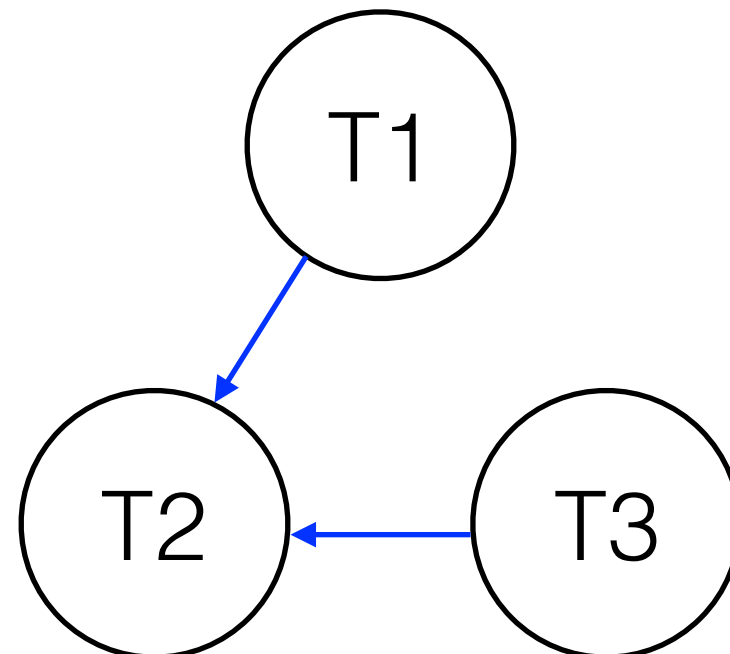
Worksheet - 2a

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

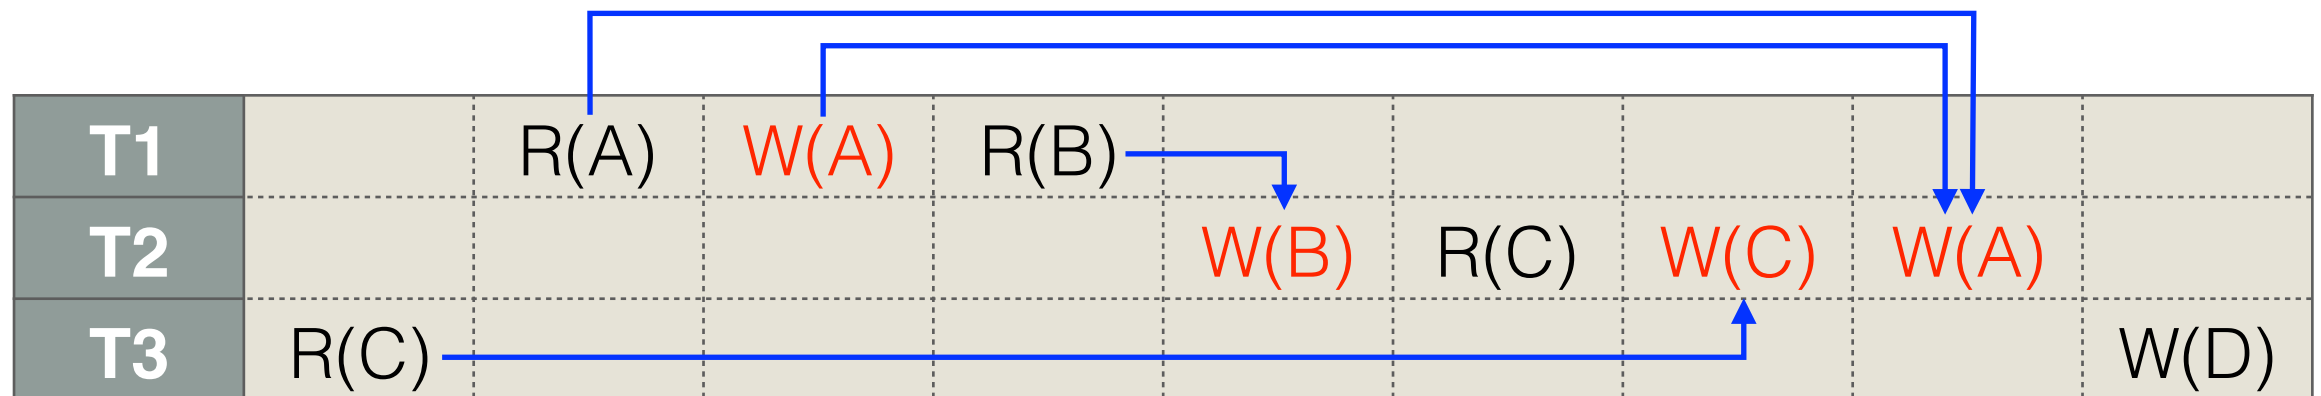


Blue arrows indicating dependencies: T1's R(A) depends on T3's R(C); T1's W(A) depends on T1's R(A); T1's R(B) depends on T2's W(B); T1's W(A) depends on T2's W(A); T2's W(C) depends on T3's R(C).

- Dependency Graph
 - Conflicting writes? Edges?

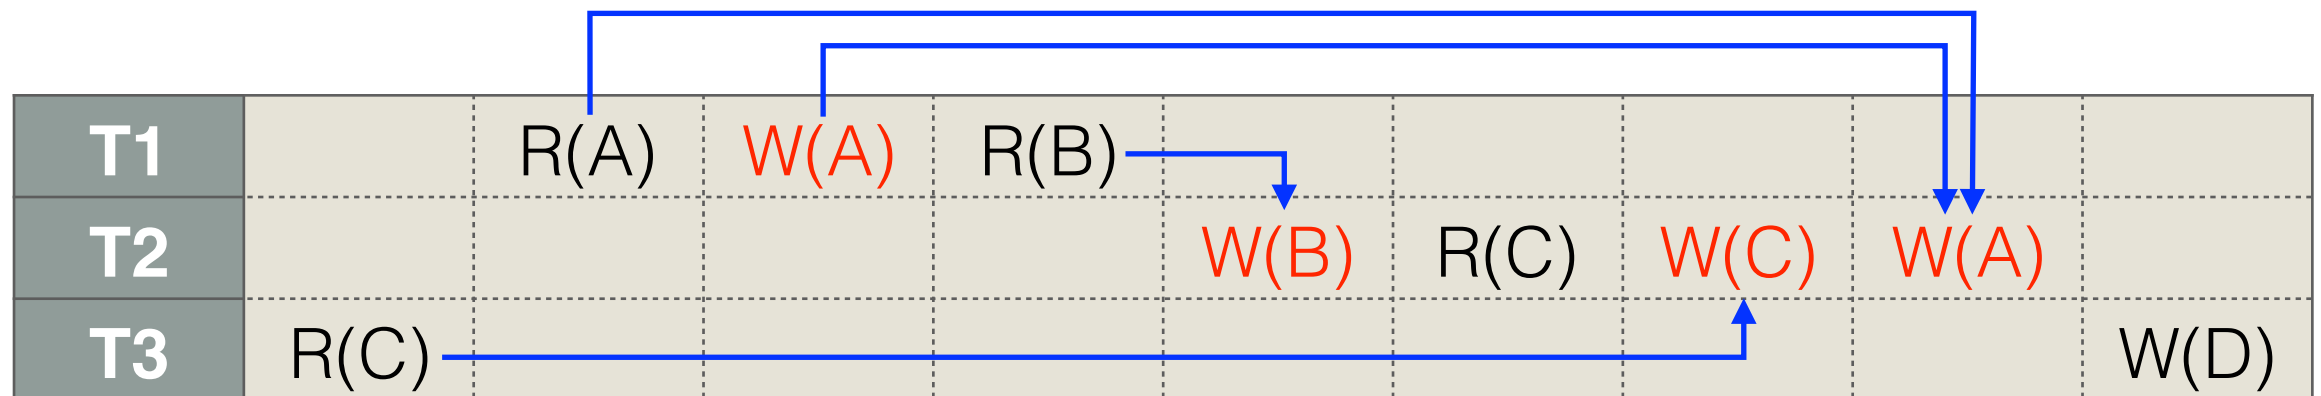


Worksheet - 2b



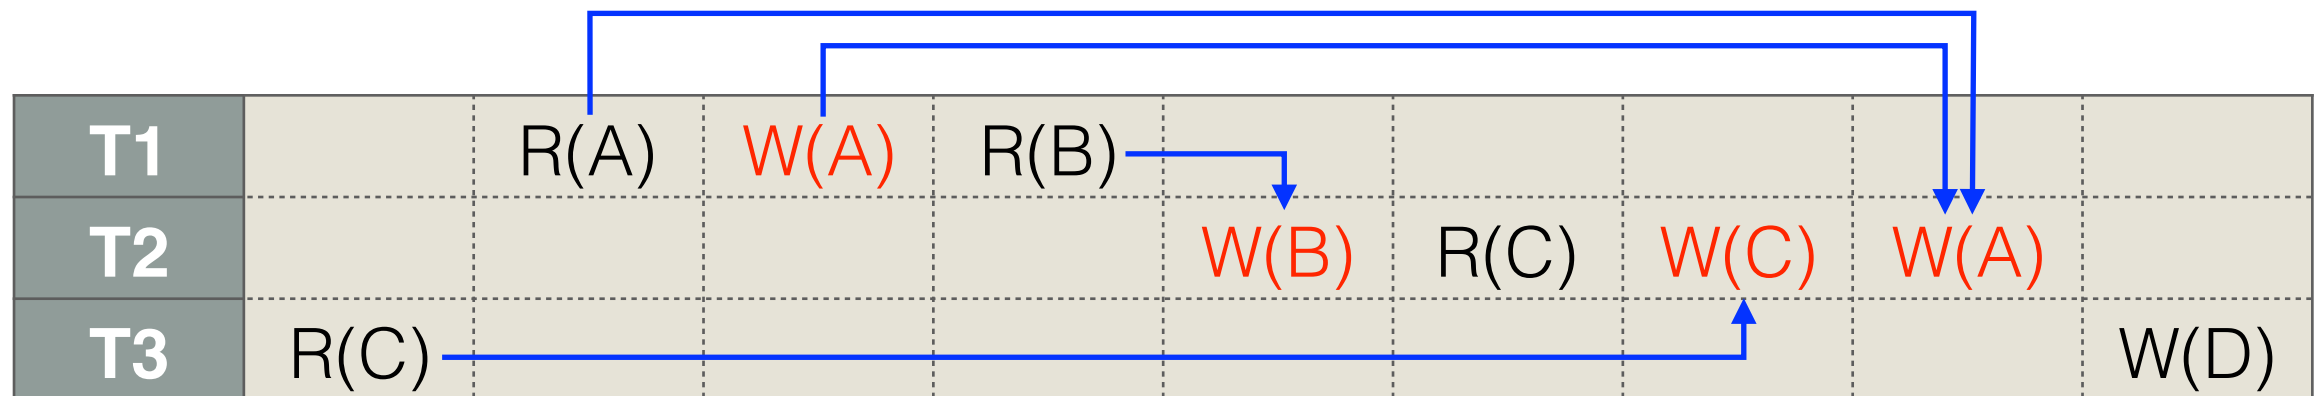
- Conflict Serializability
 - Dependency Graph?
 - What actions would we move?

Worksheet - 2b



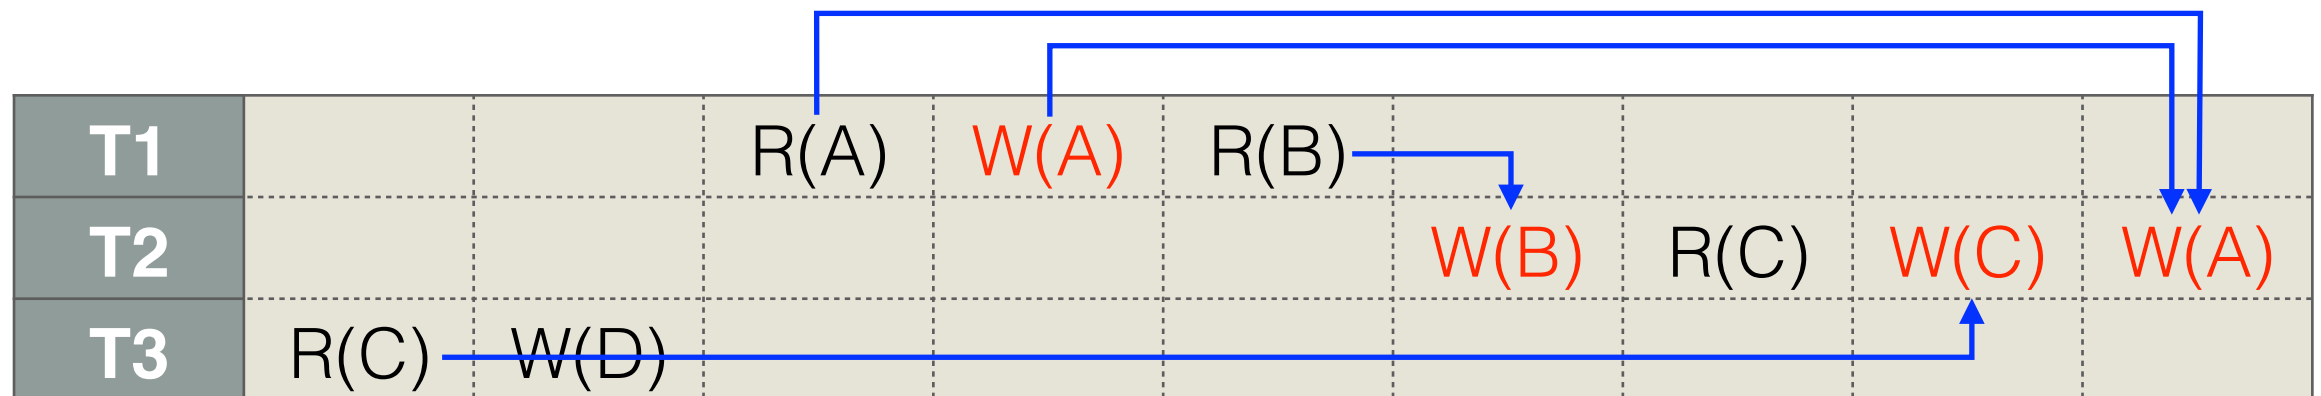
- Conflict Serializability
 - Dependency Graph? *Acyclic!*
 - What actions would we move?

Worksheet - 2b



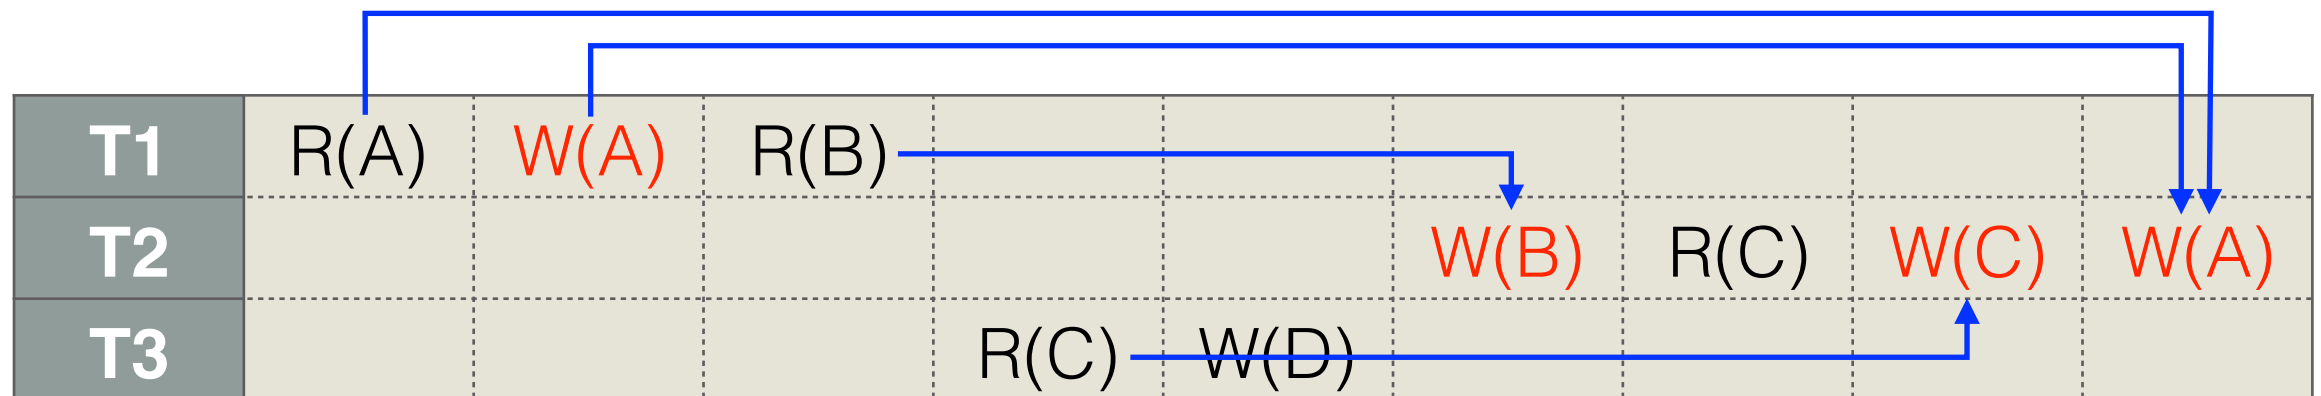
- Conflict Serializability
 - Dependency Graph? **Acyclic!**
 - What actions would we move? **T3**

Worksheet - 2b



- Conflict Serializability
 - Dependency Graph? **Acyclic!**
 - What actions would we move? **T3**

Worksheet - 2b



- Conflict Serializability
 - Dependency Graph? **Acyclic!**
 - What actions would we move? **T3**

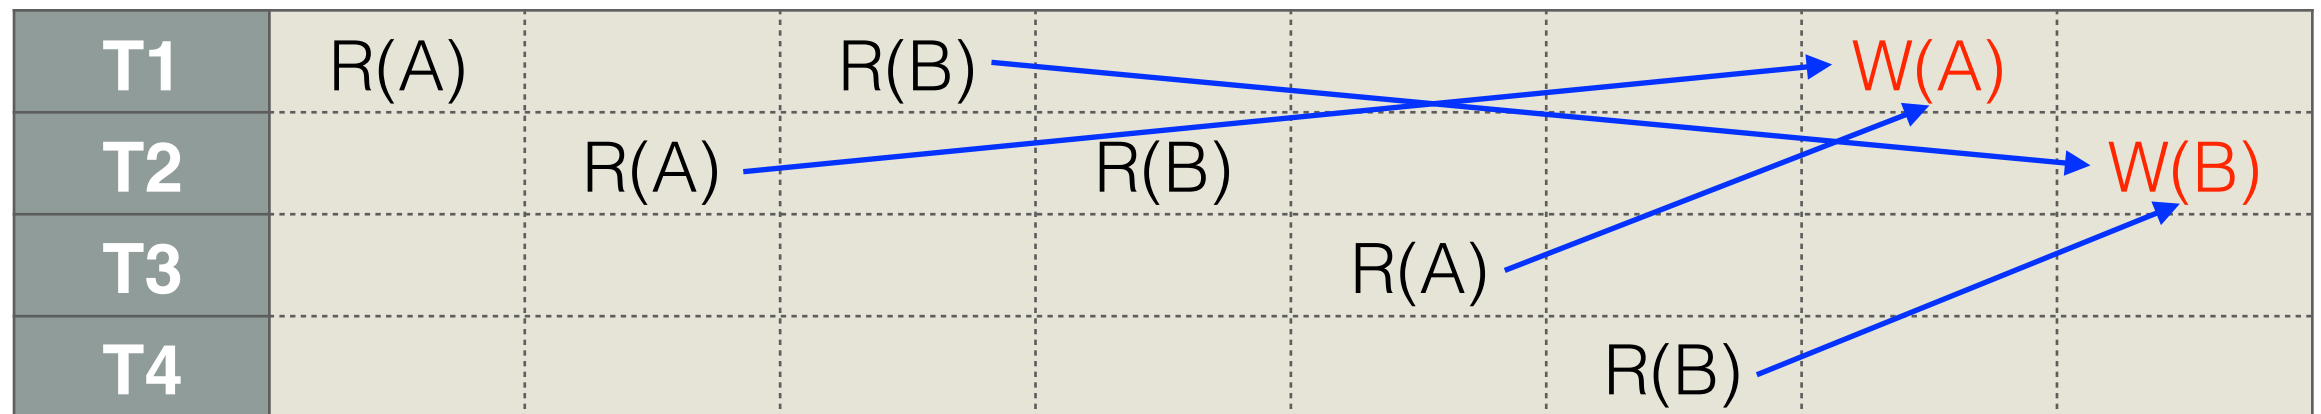
Worksheet - 2c

T1	R(A)		R(B)				W(A)	
T2		R(A)		R(B)				W(B)
T3					R(A)			
T4						R(B)		

Worksheet - 2c

T1	R(A)		R(B)				W(A)	
T2		R(A)		R(B)				W(B)
T3					R(A)			
T4						R(B)		

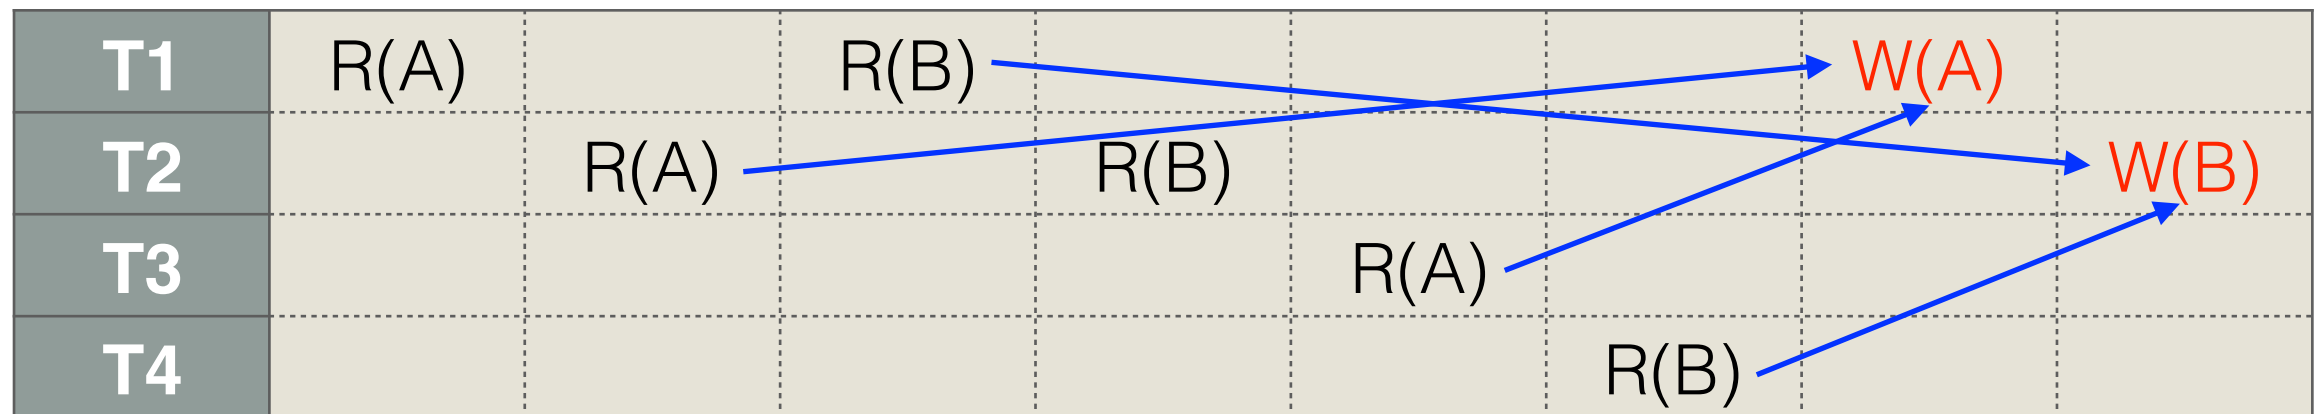
Worksheet - 2c



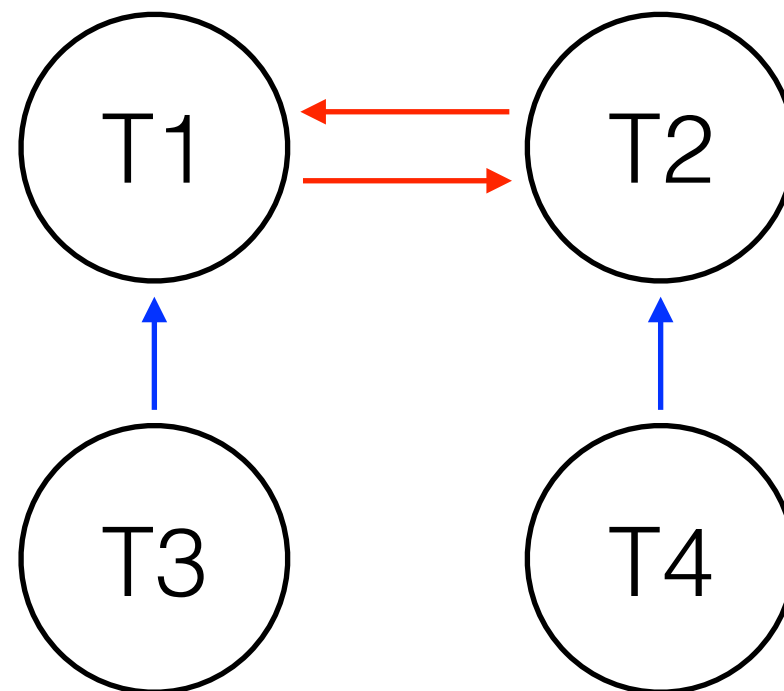
- Is this conflict serializable?

Worksheet - 2c

T1	R(A)		R(B)				W(A)	
T2		R(A)	R(B)					W(B)
T3					R(A)			
T4						R(B)		



- Is this conflict serializable?



2PL

- S = shared lock
 - reads
- X = exclusive lock
 - writes (and reads)

	S	X
S	*	
X		

- For each xact: Once a lock has been released, cannot acquire any more

What does 2PL guarantee?

- Conflict Serializability?
- Cascading aborts?



What does 2PL guarantee?

- Conflict Serializability?



- Cascading aborts?

What does 2PL guarantee?

- Conflict Serializability? 
- Cascading aborts? 

Strict 2PL

- Like 2PL, but addresses cascading aborts
 - Release locks only when a xact completes
 - Completion = commit, or abort + rollback




Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	
B = B*10	
Write(B)	
Lock_X(F)	
Unlock(B)	
	Lock_S(F)
F = B*100	
Write(F)	
Unlock(F)	
	Read(F)
	Unlock(F)
	Lock_S(B)
	Read(B)
	Print(F+B)
	Unlock(B)

- 2PL?
- Strict 2PL?
- No deadlock?

Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	
B = B*10	
Write(B)	
Lock_X(F)	
Unlock(B)	
	Lock_S(F)
F = B*100	
Write(F)	
Unlock(F)	
	Read(F)
	Unlock(F)
	Lock_S(B)
	Read(B)
	Print(F+B)
	Unlock(B)

- 2PL? 
- Strict 2PL? 
- No deadlock? 

Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	
B = B*10	
Write(B)	
Lock_X(F)	
Unlock(B)	
	Lock_S(F)
F = B*100	
Write(F)	
Unlock(F)	
	Read(F)
	Unlock(F)
	Lock_S(B)
	Read(B)
	Print(F+B)
	Unlock(B)

- Conflict serializable?
- View Serializable?
- Serializable?

Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	
B = B*10	
Write(B)	
Lock_X(F)	
Unlock(B)	
	Lock_S(F)
F = B*100	
Write(F)	
Unlock(F)	
	Read(F)
	Unlock(F)
	Lock_S(B)
	Read(B)
	Print(F+B)
	Unlock(B)

- Conflict serializable? ✓

- View Serializable? ✓

- Serializable? ✓

Deadlock Detection

- “Waits-for” Graph:
 - edge from T_i to T_j if i is waiting for j
- Deadlock if cycle in graph

Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

- What needs to wait?

T1

T2

T3

T4

Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

- What needs to wait?

T1 → T2

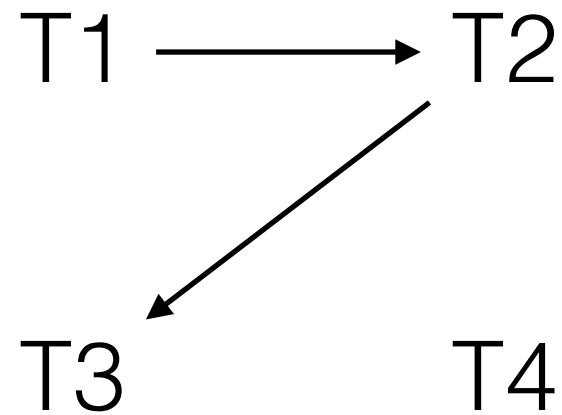
T3

T4

Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

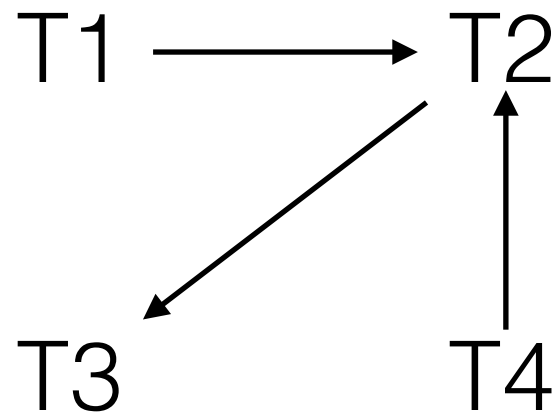
- What needs to wait?



Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

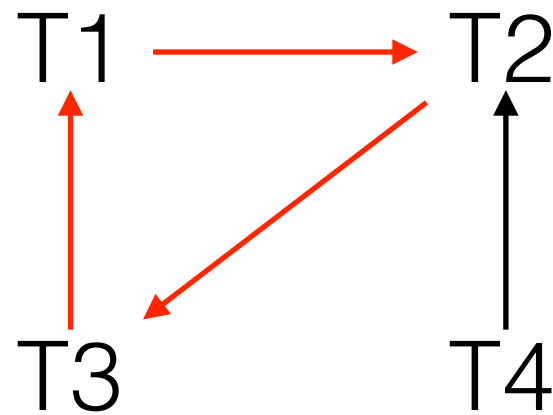
- What needs to wait?



Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

- What needs to wait?



Deadlock Avoidance

- Priorities based on timestamp
- Wait-Die: if T_i higher, wait for T_j , else *abort T_i*
- Wound-Wait: if T_i higher, *T_j aborts*, else wait for T_j

Locking Granularity

- Databases have hierarchies
 - (Top) Databases > tables > pages > tuples
 - Assign locks top down, release bottom up
- Use intent locks!
 - IS - intent to get S
 - IX - intent to get X
 - SIX - S and IX

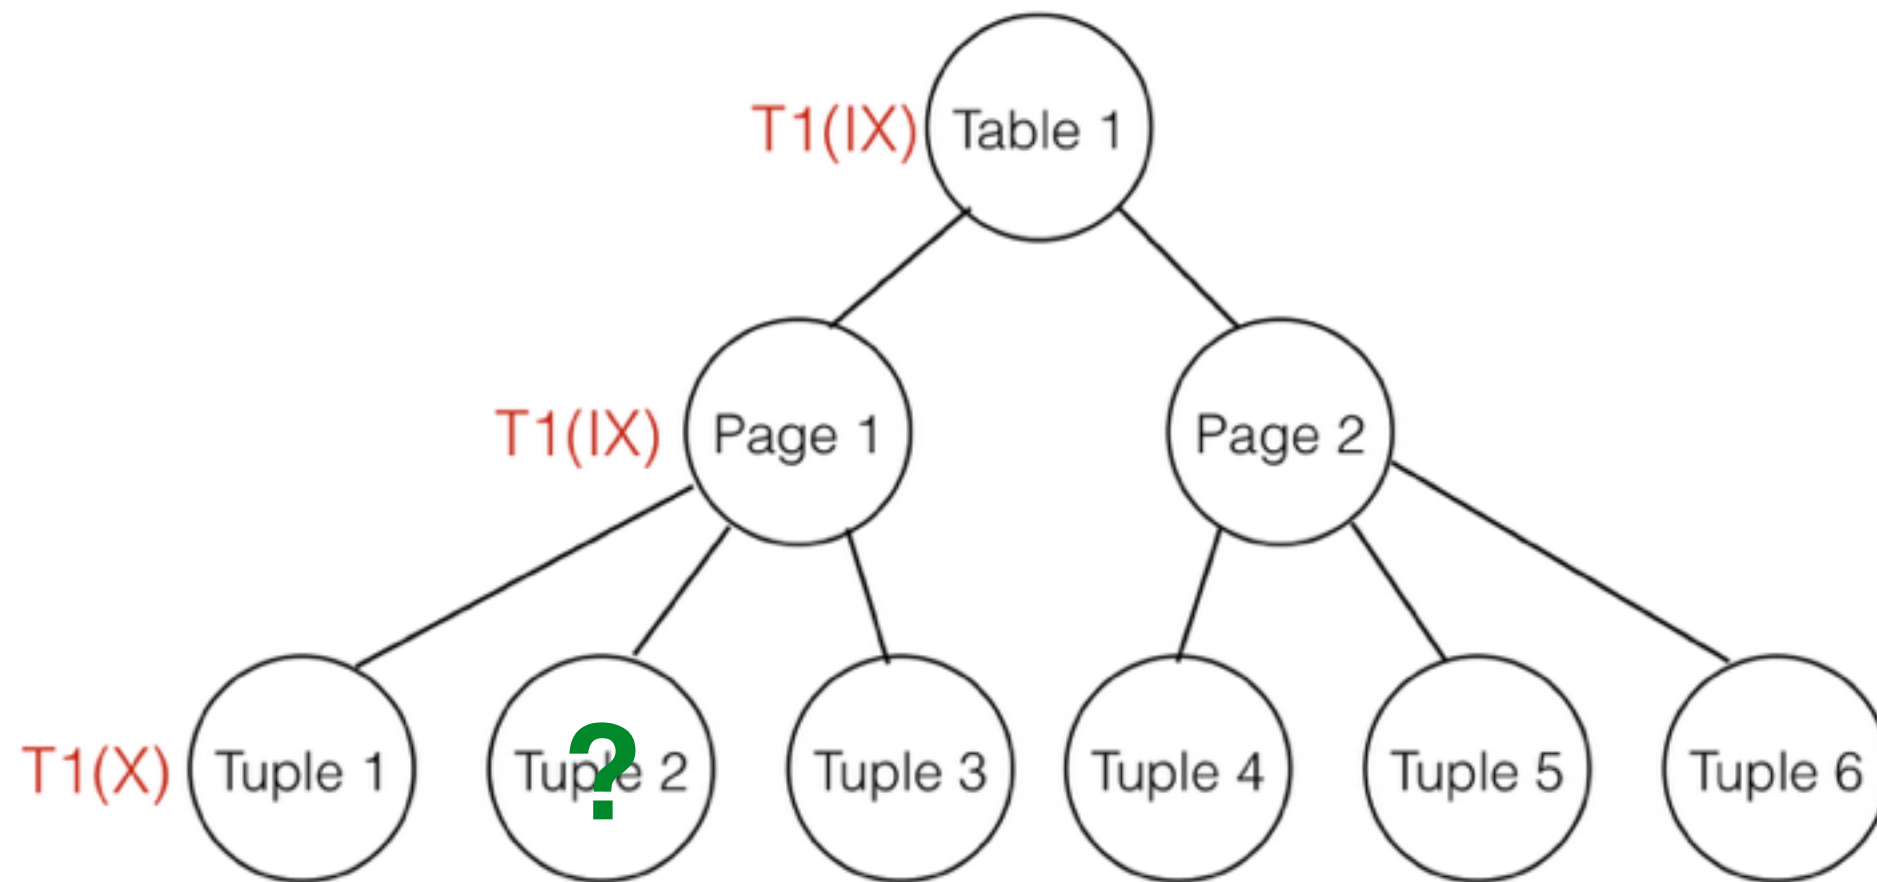
Locking Granularity

	IS	IX	SIX	S	X
IS	✓	✓	✓	✓	—
IX	✓	✓	—	—	—
SIX	✓	—	—	—	—
S	✓	—	—	✓	—
X	—	—	—	—	—

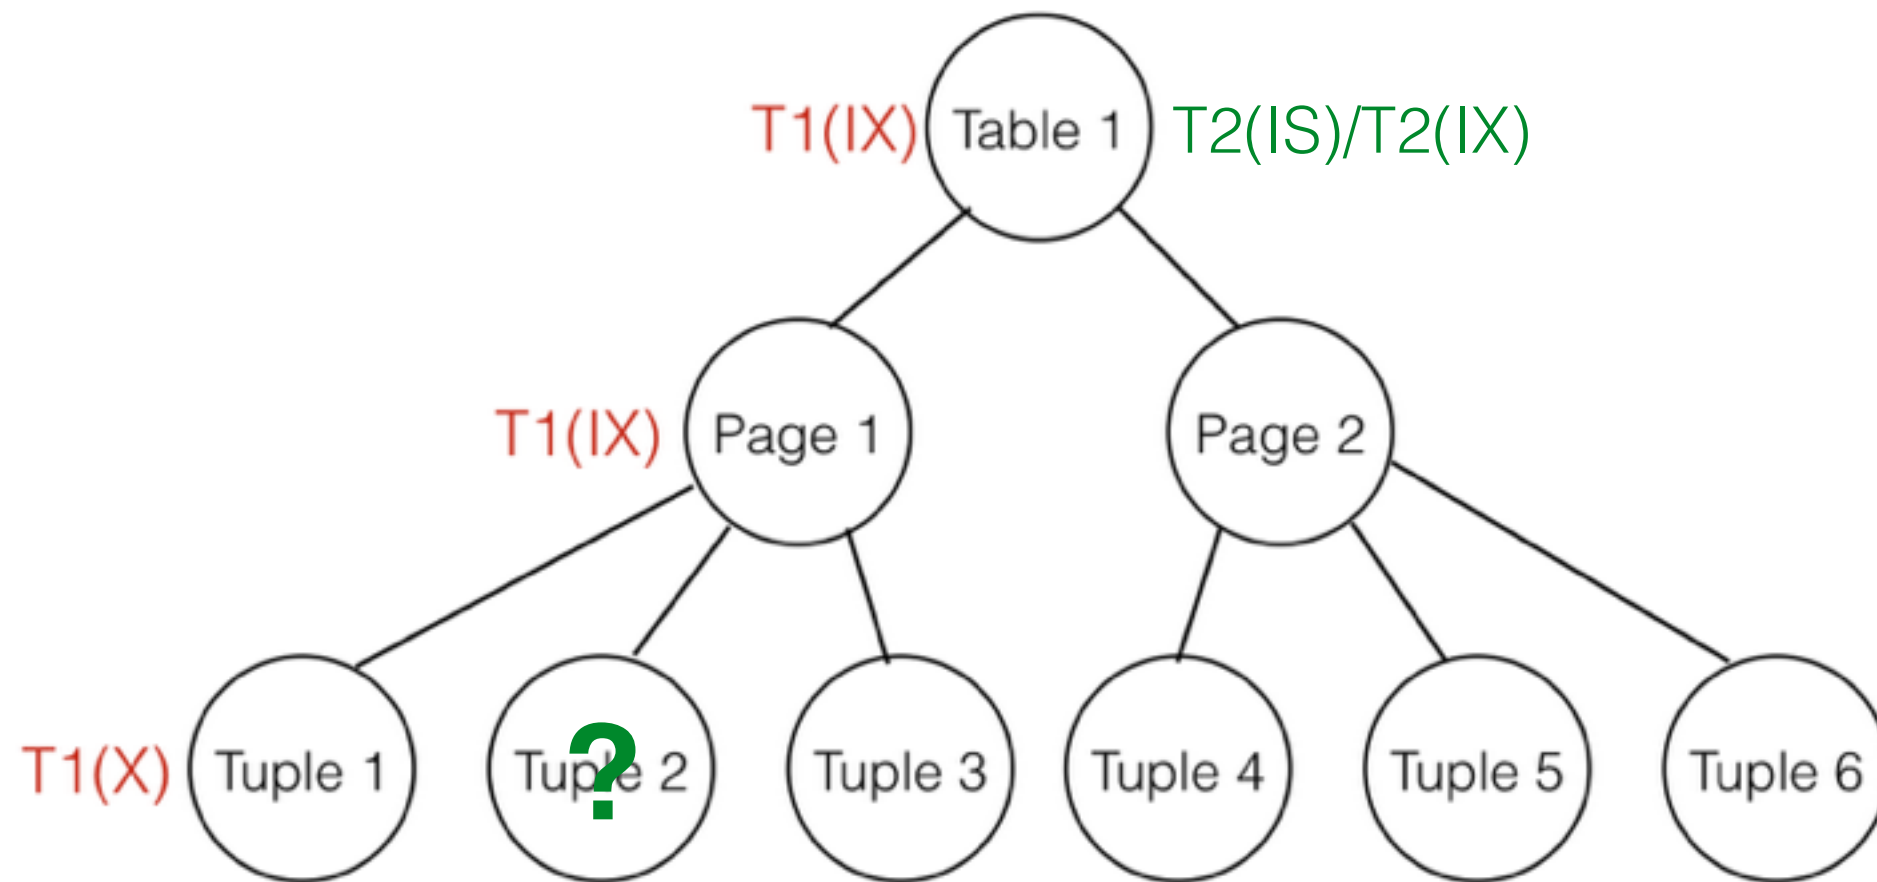
From Wikipedia:

To Get	Must Have on all Ancestors
IS or S	IS or IX
IX, SIX or X	IX or SIX

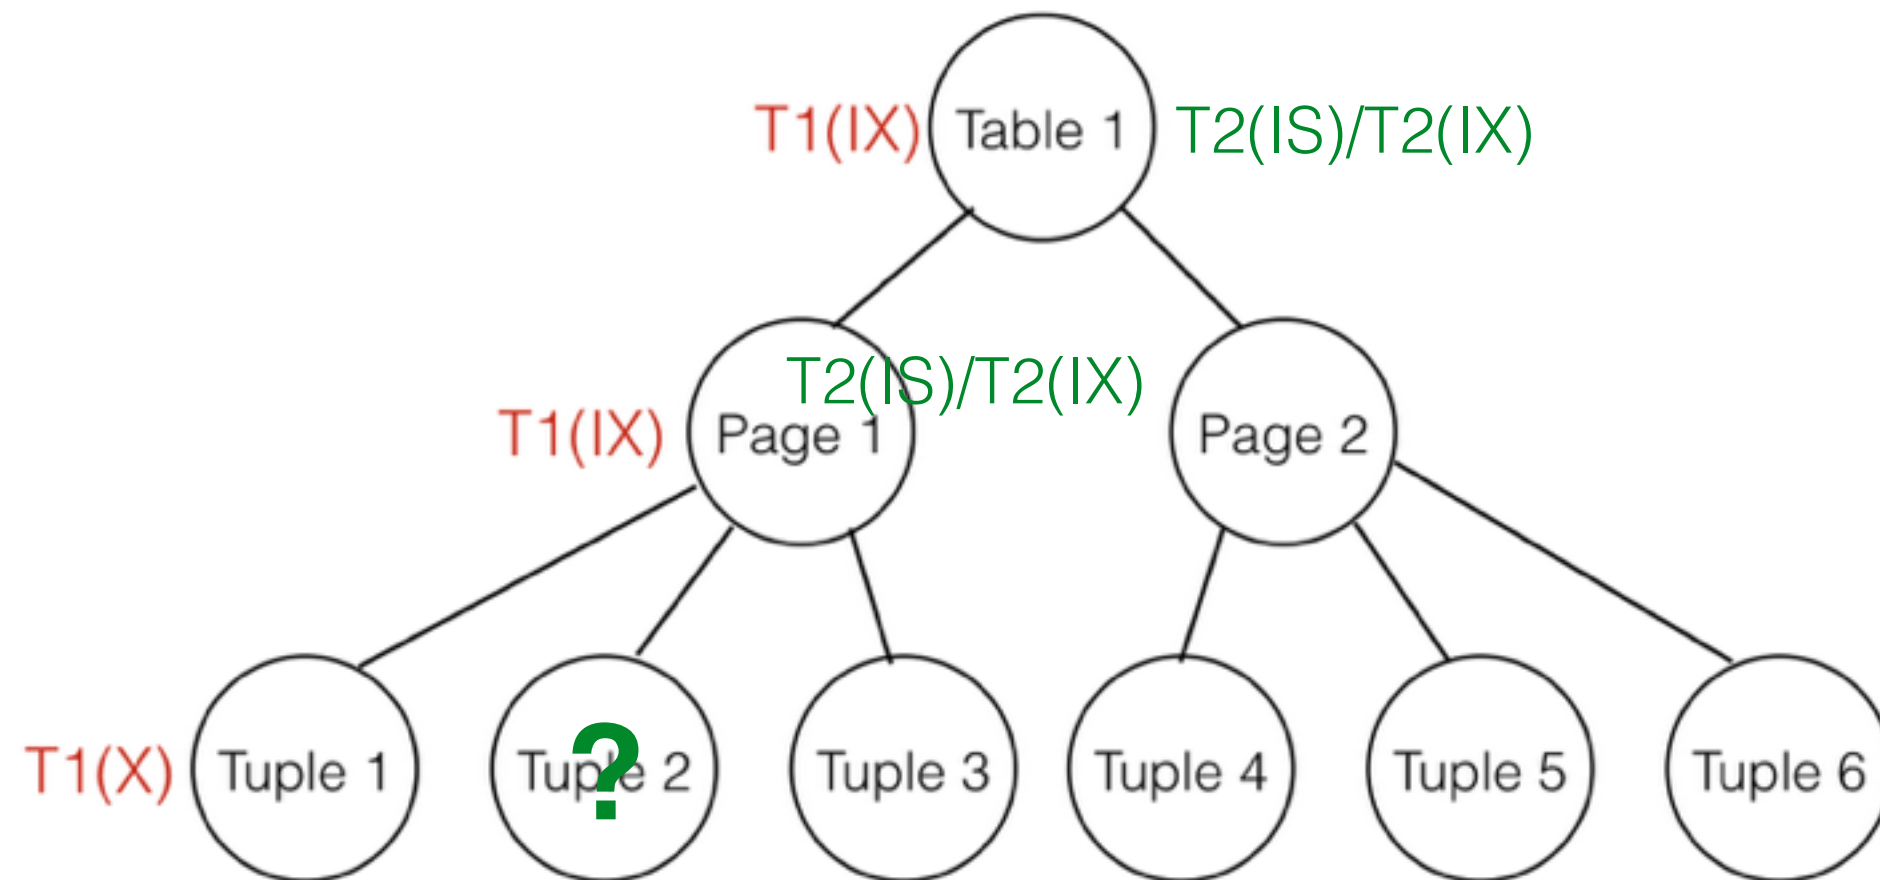
Worksheet - #3



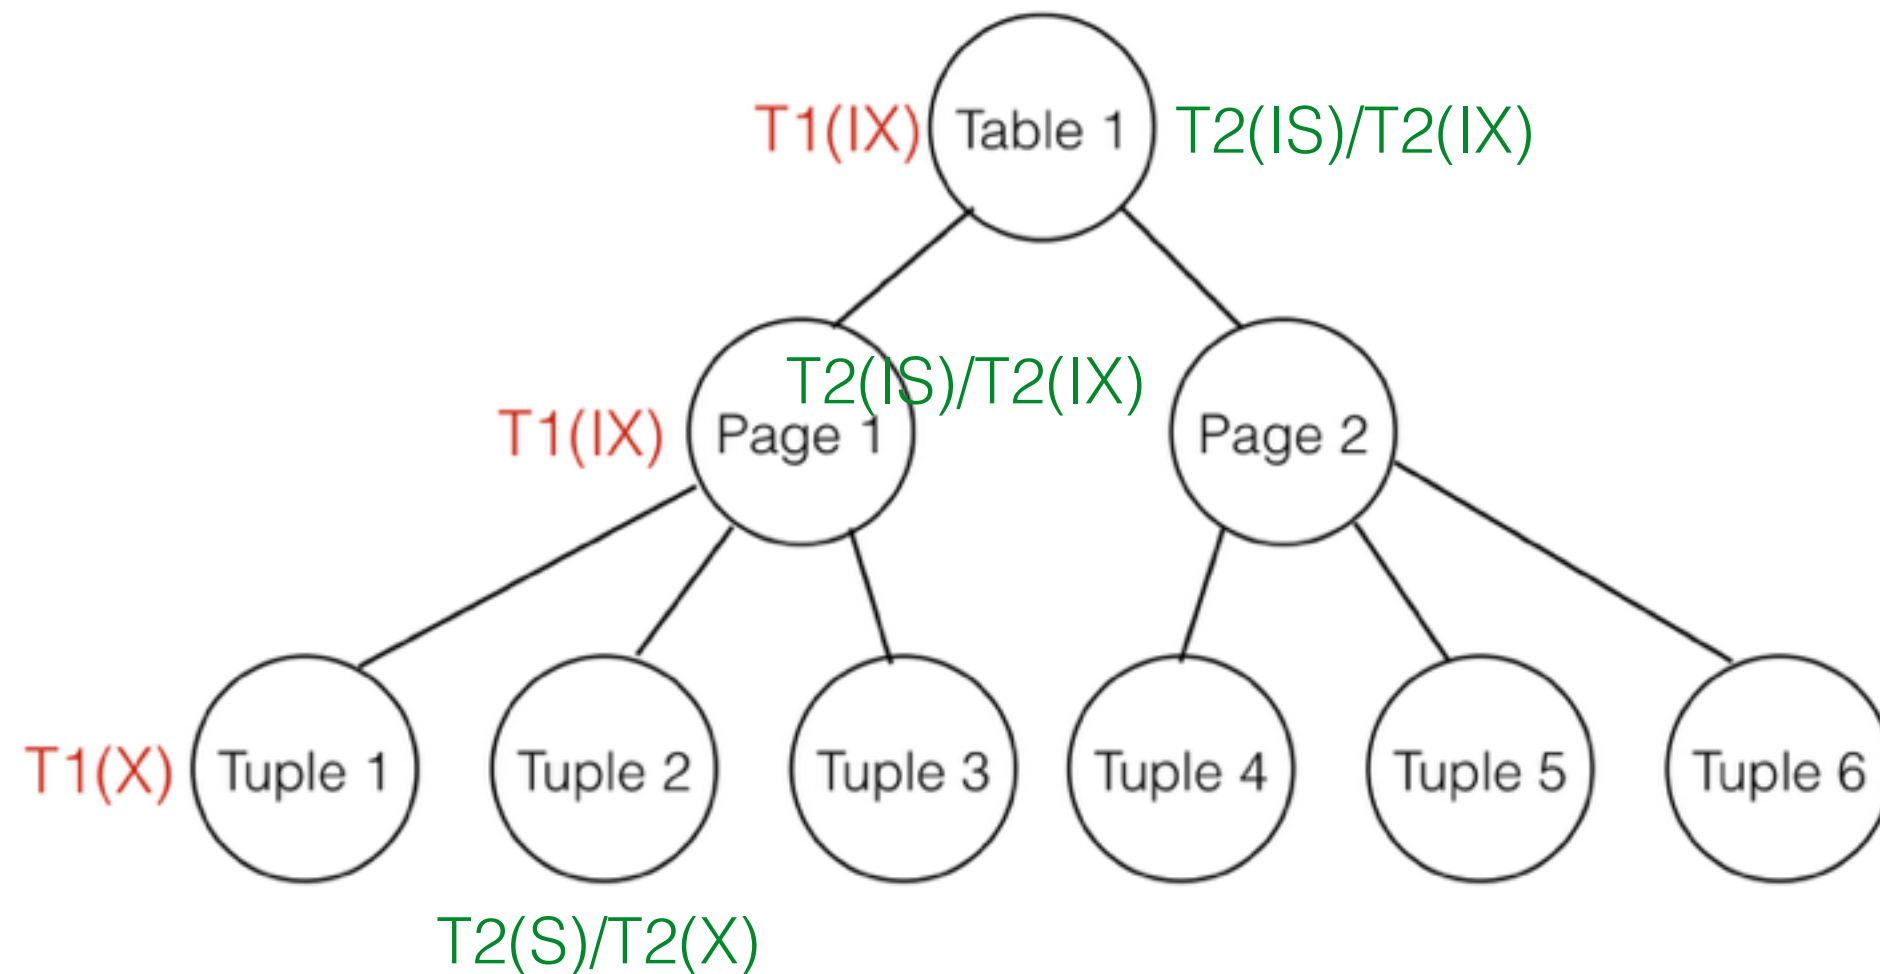
Worksheet - #3



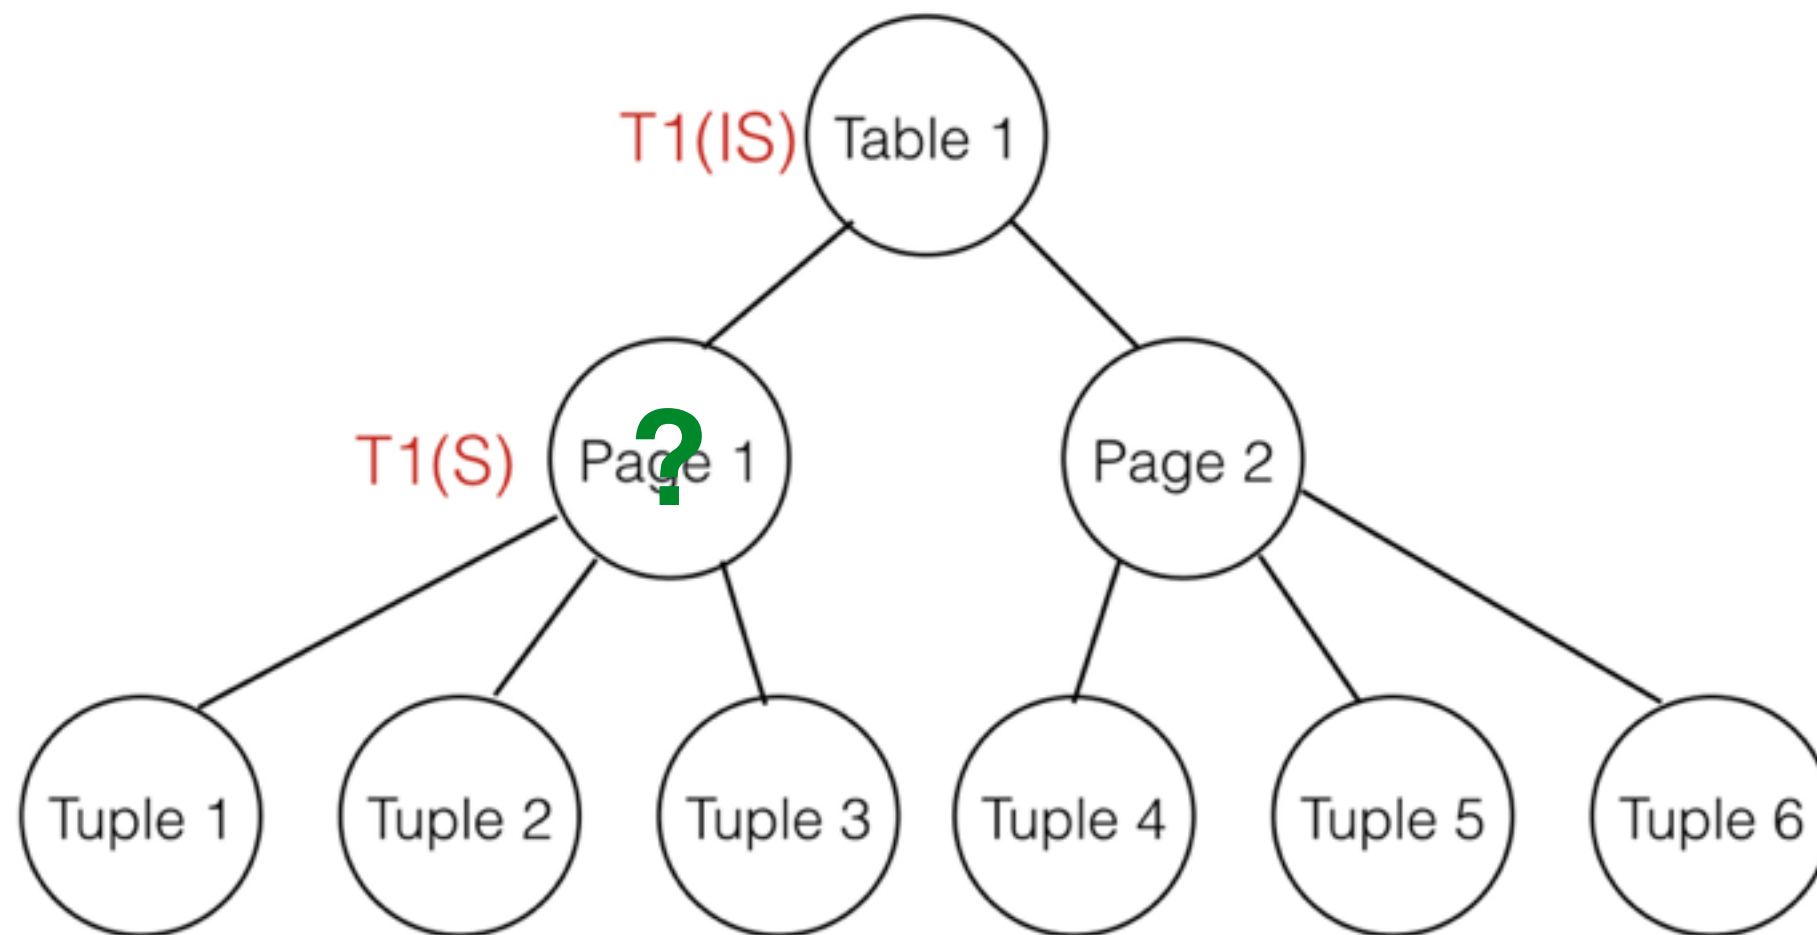
Worksheet - #3



Worksheet - #3



Worksheet - #3



Worksheet - #3

