

CS 186 Discussion #8

Transactions, Concurrency Control

Logistics

- Homework 4 due 11/04
- Midterm 2 on 11/09
 - Up to concurrency control (lock granularity)
- Midterm 3 partially cumulative

System R...

```
SELECT *  
FROM R, S, T  
WHERE R.a = S.a  
AND S.b = T.b;
```

We now add the third table and have the following join costs:

- 1) (R join S) join T = 10,000
- 2) T join (R join S) = 6,000
- 3) (S join R) join T = 15,000
- 4) T join (S join R) = 11,000
- 5) (R join T) join S = 10,000
- 6) S join (R join T) = 7,000
- 7) (T join R) join S = 14,000
- 8) S join (T join R) = 16,000
- 9) (S join T) join R = 13,000
- 10) R join (S join T) = 12,000
- 11) (T join S) join R = 20,000
- 12) R join (T join S) = 9,000

These are the two-table join costs:

- 1) R join S = 6,000
- 2) S join R = 2,000
- 3) R join T = 5,000
- 4) T join R = 1,000
- 5) S join T = 4,000
- 6) T join S = 3,000

Which of these will the optimizer select as your final query plan? *

Which of the following two-table join plans will be selected by a System R-style query optimizer? *

System R...

```
SELECT *  
FROM R, S, T  
WHERE R.a = S.a  
AND S.b = T.b;
```

We now add the third table and have the following join costs:

- 1) (R join S) join T = 10,000
- 2) T join (R join S) = 6,000
- 3) (S join R) join T = 15,000 ←
- 4) T join (S join R) = 11,000
- 5) (R join T) join S = 10,000
- 6) S join (R join T) = 7,000
- 7) (T join R) join S = 14,000
- 8) S join (T join R) = 16,000
- 9) (S join T) join R = 13,000
- 10) R join (S join T) = 12,000
- 11) (T join S) join R = 20,000 ←
- 12) R join (T join S) = 9,000

These are the two-table join costs:

- 1) R join S = 6,000
- 2) S join R = 2,000 ←
- 3) R join T = 5,000
- 4) T join R = 1,000
- 5) S join T = 4,000
- 6) T join S = 3,000 ←

Which of these will the optimizer select as your final query plan? *

Which of the following two-table join plans will be selected by a System R-style query optimizer? *

Transactions

- **A**tomicity: All actions in the xact happen, or none
 - Logging
- **C**onsistency: xact will not break DB consistency
 - Integrity constraints
- **I**solation: Execution of the xact is isolated from other xacts
 - Serial ordering
- **D**urability: Committed xacts have persistent effects
 - Logging

Transactions

- **A**tomicity: All actions in the xact happen, or none
 - Logging
- **C**onsistency: xact will not break DB consistency
 - Integrity constraints
- **I**solation: Execution of the xact is isolated from other xacts
 - Serial ordering
- **D**urability: Committed xacts have persistent effects
 - Logging

Transactions

- **A**tomicity: All actions in the xact happen, or none
 - Logging
- **C**onsistency: xact will not break DB consistency
 - Integrity constraints
- **I**solation: Execution of the xact is isolated from other xacts
 - Serial ordering
- **D**urability: Committed xacts have persistent effects
 - Logging

Serializability

- Serial Schedule - no intermittent transactions

✓ R(A) W(A) R(B) W(B) R(A) W(A) R(B) W(B)

✗ R(A) W(A) R(B) W(B)
R(A) W(A) R(B) W(B)

Serializability

- Serializable: equivalent to any serial schedule
- Equivalent: same xacts/actions and final state

R(A) W(A)

 R(B) W(B)

Equivalent?

R(A) W(A)

 R(B) W(B)

Serializable?

Serializability

- Serializable: equivalent to any serial schedule
- Equivalent: same xacts/actions and final state

R(A) W(A)

 R(B) W(B)

Equivalent?



R(A) W(A)

 R(B) W(B)

Serializable?



Conflict Serializability

- What is a conflict?
 - Different xacts on one object that is *written to*

T1: R(A) R(B) W(A)

T2: R(B) W(B)

Conflict Serializability

- What is a conflict?
 - Different xacts on one object that is *written to*


T1: R(A) R(B) W(A)

T2: R(B) W(B)

Conflict Serializability

T1: R(A) R(B) W(A)


T2: R(B) W(B)



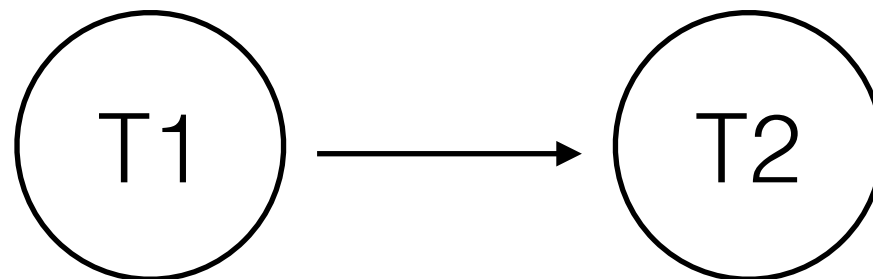
- Dependency Graph:
 - One node per xact
 - Edge from T_i to T_j if some operation O_i is earlier than and conflicts with O_j

Conflict Serializability

T1: R(A) R(B) W(A)
T2: R(B) W(B)



- Dependency Graph:
 - One node per xact
 - Edge from T_i to T_j if some operation O_i is earlier than and conflicts with O_j



Conflict Serializability

- Conflict Serializable: conflict equivalent to some serial schedule
 - Conflict Equivalent: same xacts/actions and *every conflict pair is ordered the same way*
- Serializable vs. Conflict Serializable?

Conflict Serializability

- Conflict Serializable: conflict equivalent to some serial schedule
 - Conflict Equivalent: same xacts/actions and *every conflict pair is ordered the same way*
 - Serializable vs. Conflict Serializable?
- **Theorem**: Schedule is conflict serializable *if and only if* dependency graph is acyclic

Worksheet - 2a

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

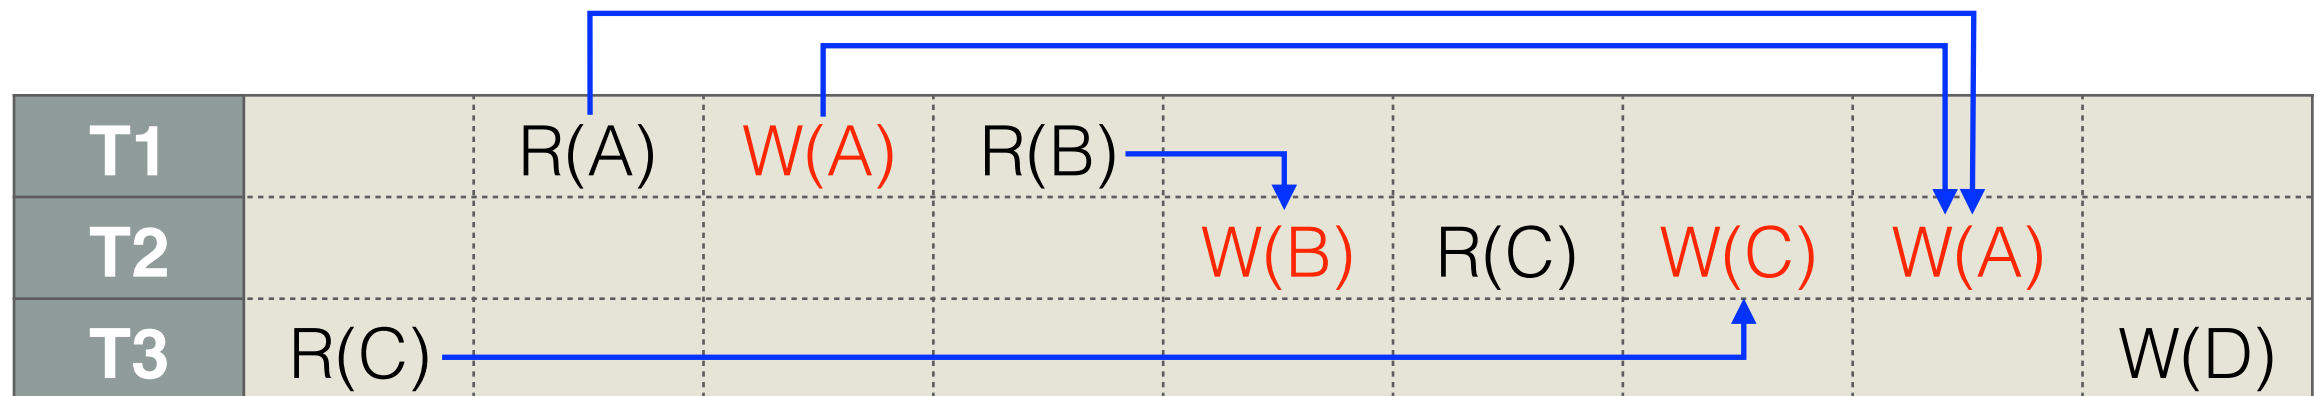
- Dependency Graph
 - Conflicting writes? Edges?

Worksheet - 2a

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

- Dependency Graph
 - Conflicting writes? Edges?

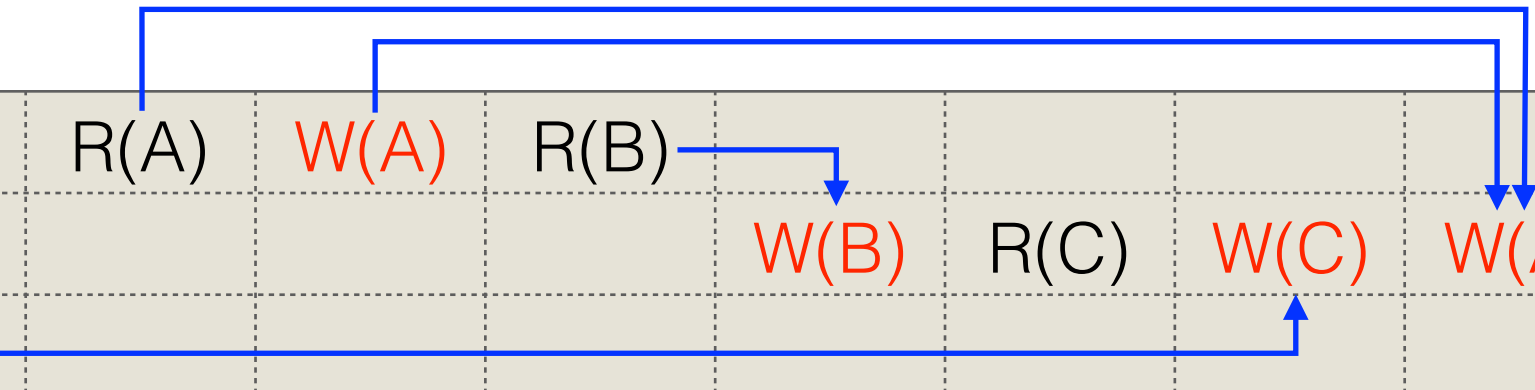
Worksheet - 2a



- Dependency Graph
 - Conflicting writes? Edges?

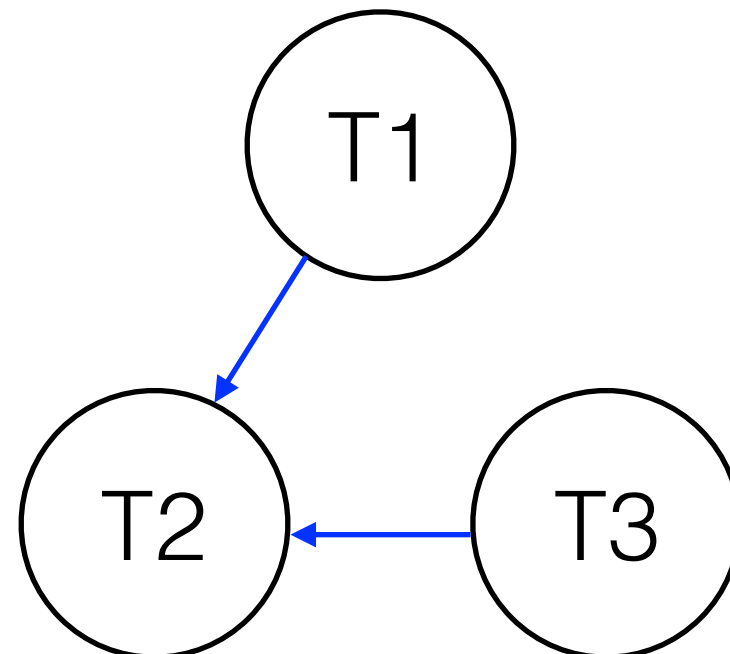
Worksheet - 2a

T1		R(A)	W(A)	R(B)					
T2					W(B)	R(C)	W(C)	W(A)	
T3	R(C)								W(D)

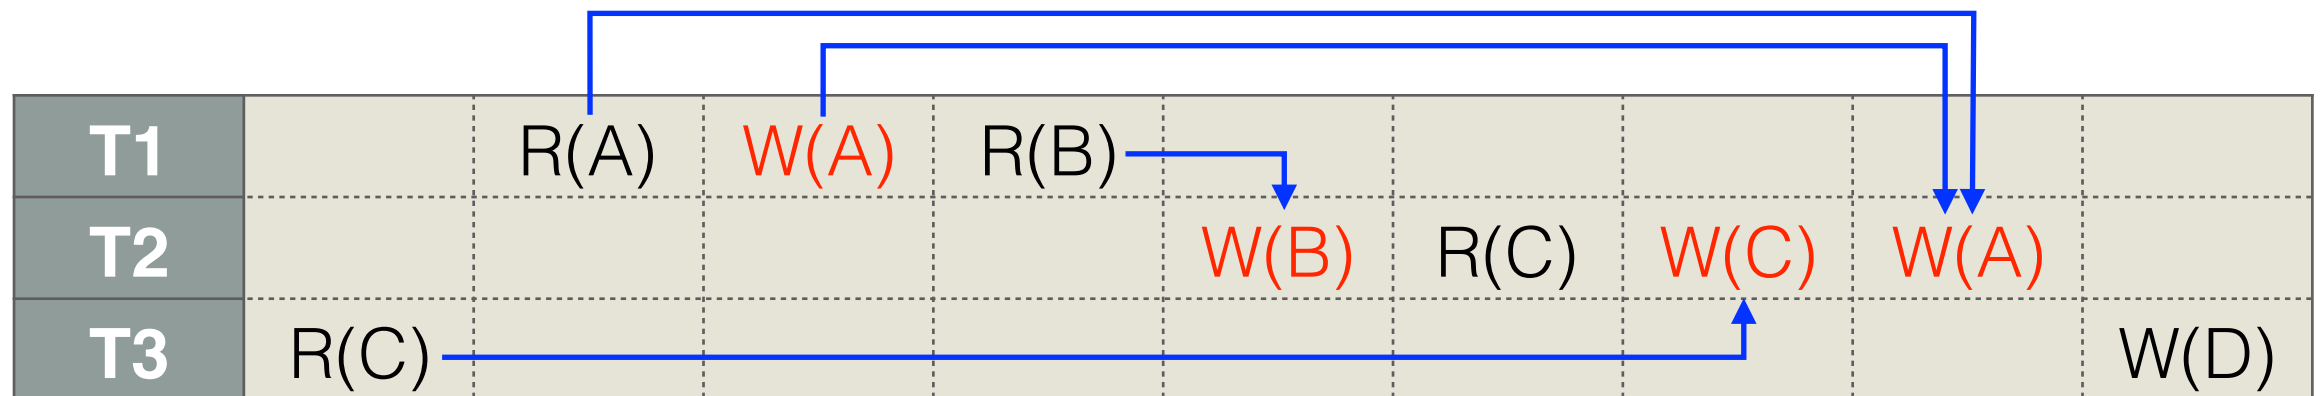


Blue arrows indicating dependencies: T1's R(A) depends on T3's R(C); T1's W(A) depends on T1's R(A); T1's R(B) depends on T2's W(B); T1's W(A) depends on T2's W(A); T2's W(C) depends on T3's R(C).

- Dependency Graph
 - Conflicting writes? Edges?

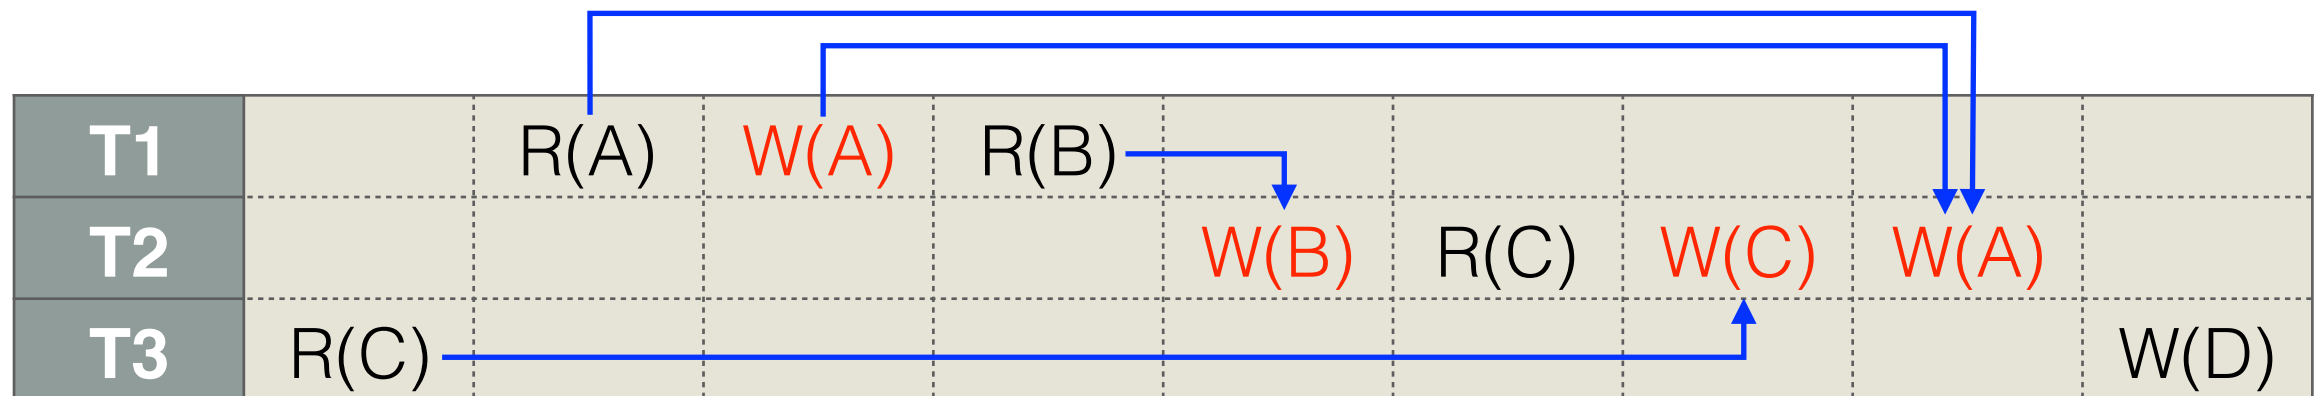


Worksheet - 2b



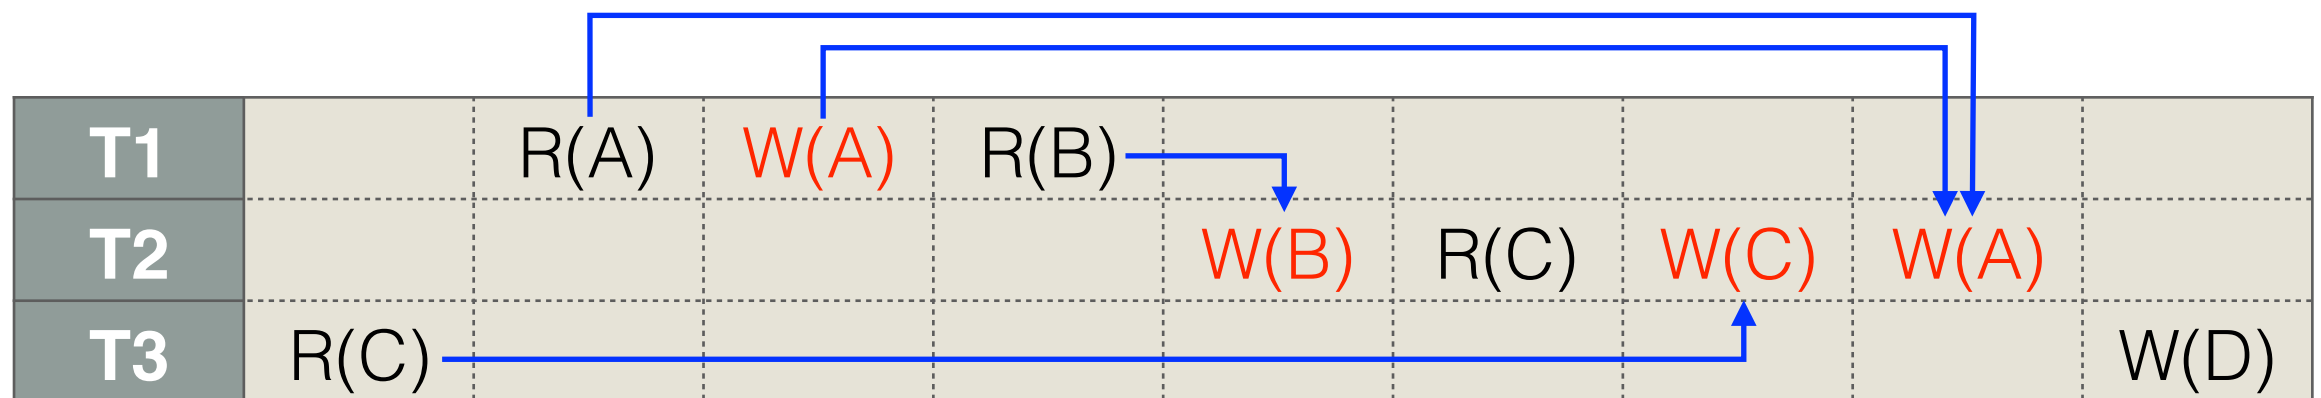
- Conflict Serializability
 - Dependency Graph?
 - What actions would we move?

Worksheet - 2b



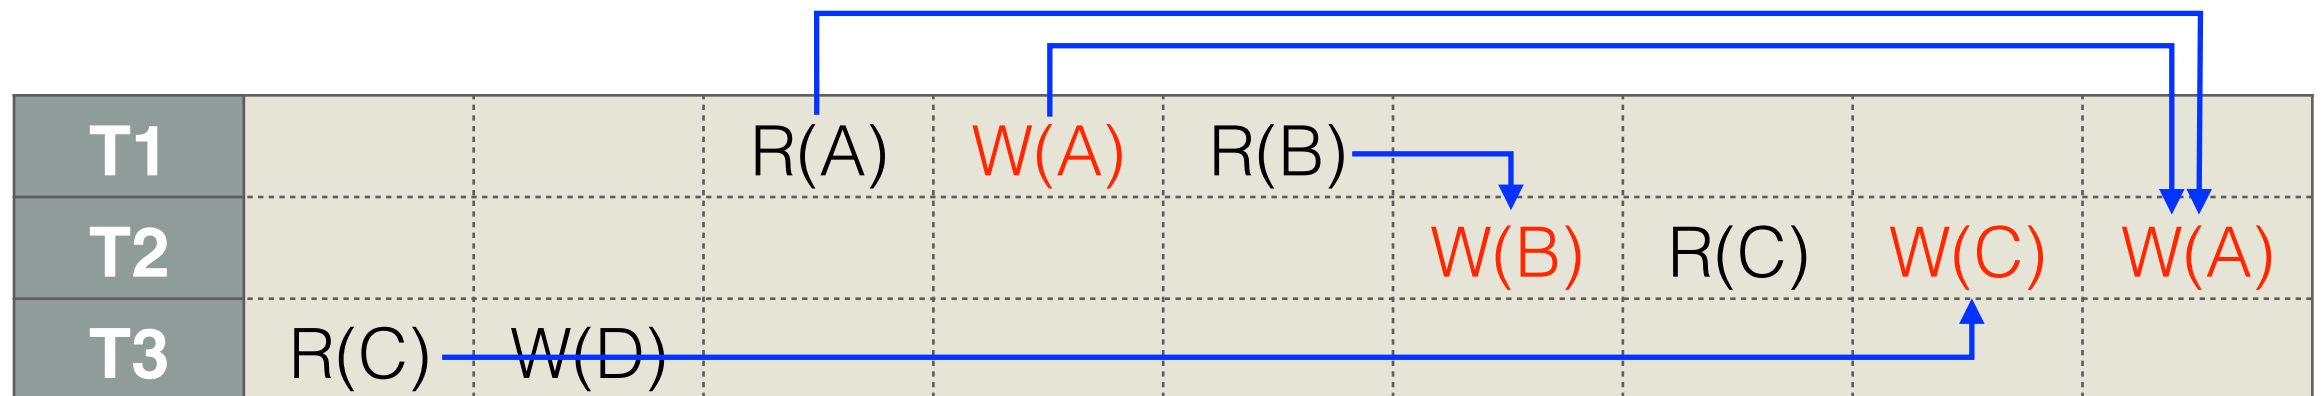
- Conflict Serializability
 - Dependency Graph? *Acyclic!*
 - What actions would we move?

Worksheet - 2b



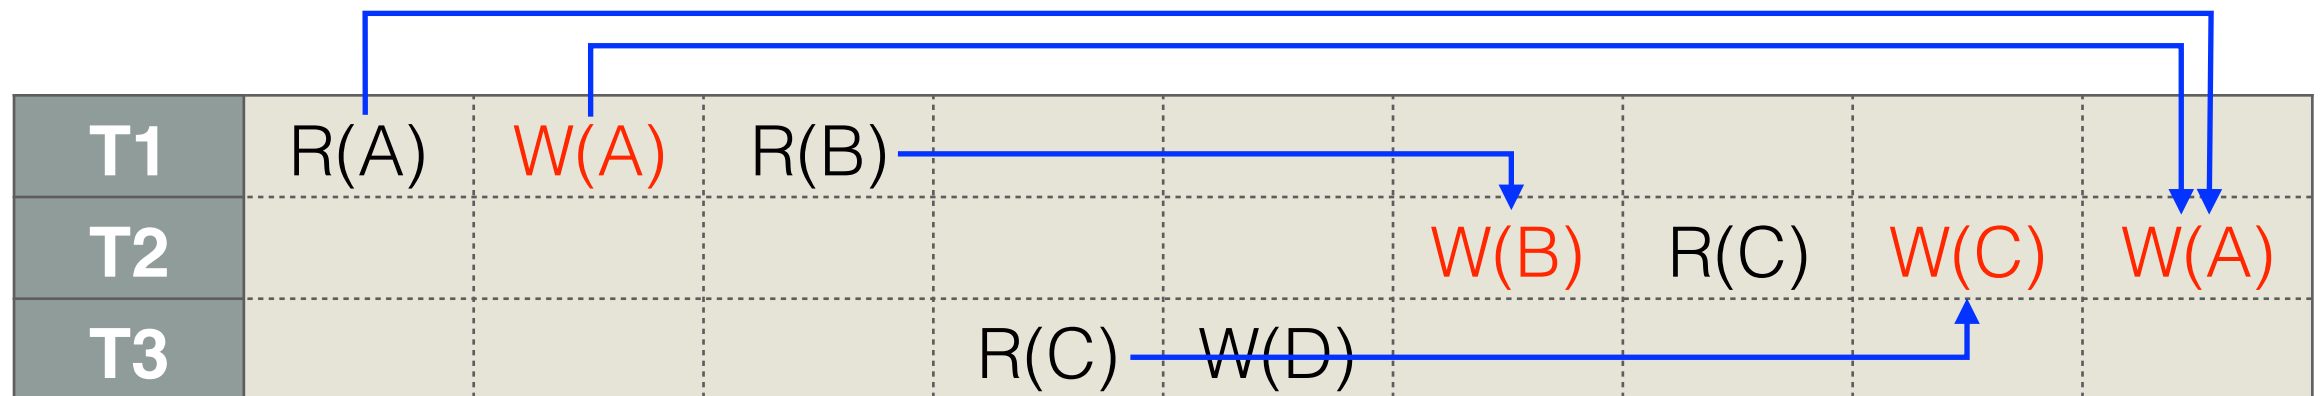
- Conflict Serializability
 - Dependency Graph? **Acyclic!**
 - What actions would we move? **T3**

Worksheet - 2b



- Conflict Serializability
 - Dependency Graph? **Acyclic!**
 - What actions would we move? **T3**

Worksheet - 2b



- Conflict Serializability
 - Dependency Graph? *Acyclic!*
 - What actions would we move? *T3*

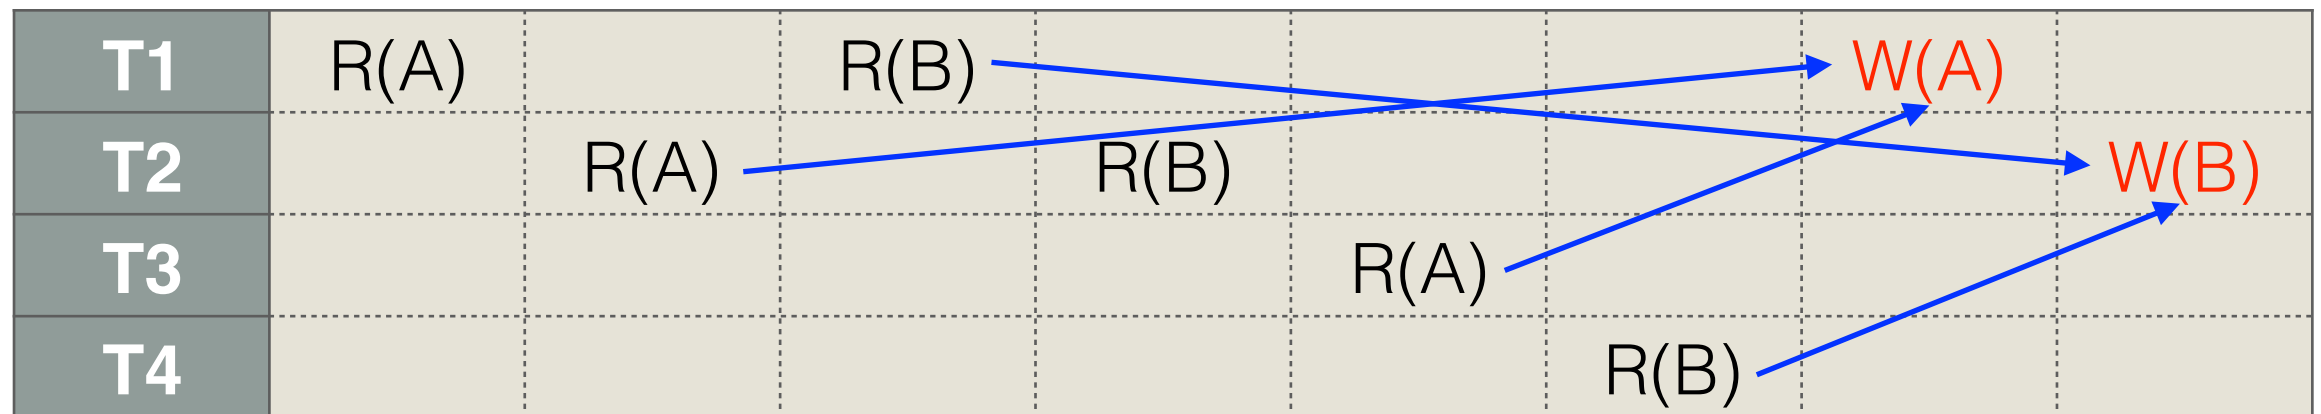
Worksheet - 2c

T1	R(A)		R(B)				W(A)	
T2		R(A)		R(B)				W(B)
T3					R(A)			
T4						R(B)		

Worksheet - 2c

T1	R(A)		R(B)				W(A)	
T2		R(A)		R(B)				W(B)
T3					R(A)			
T4						R(B)		

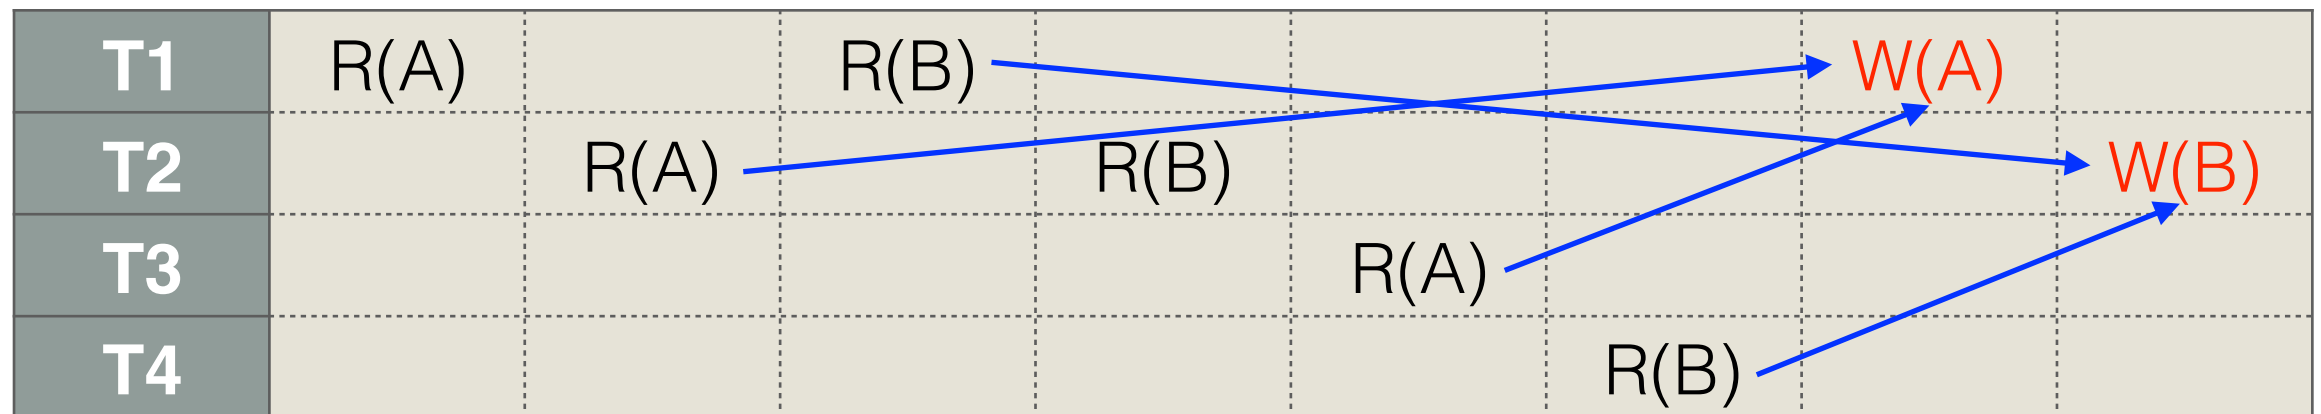
Worksheet - 2c



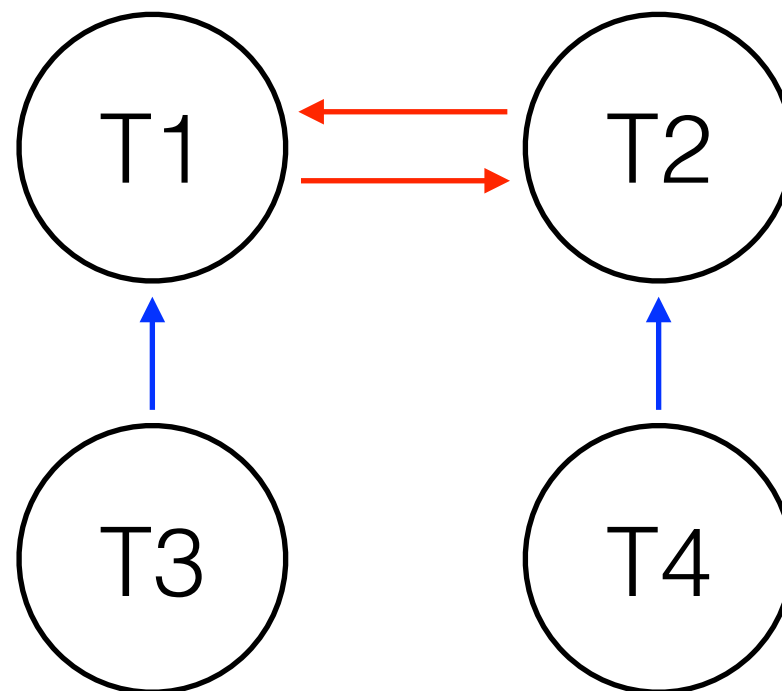
- Is this conflict serializable?

Worksheet - 2c

T1	R(A)		R(B)				W(A)	
T2		R(A)	R(B)					W(B)
T3					R(A)			
T4						R(B)		



- Is this conflict serializable?



2PL

- S = shared lock
 - reads
- X = exclusive lock
 - writes (and reads)



	S	X
S	*	
X		

- For each xact: Once a lock has been released, cannot acquire any more

What does 2PL guarantee?

- Conflict Serializability?
- Cascading aborts?

What does 2PL guarantee?

- Conflict Serializability? 
- Cascading aborts? 

Strict 2PL

- Like 2PL, but addresses cascading aborts
 - Release locks only when a xact completes
 - Completion = commit, or abort + rollback

Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	Lock_S(F)
$B = B * 10$	Read(F)
Write(B)	Unlock(F)
Lock_X(F)	Lock_S(B)
$F = B * 100$	
Write(F)	
Unlock(F)	
Unlock(B)	
	Read(B)
	Print(F+B)
	Unlock(B)

- Which locks have to wait?
- Is this conflict serializable?

Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	Lock_S(F)
$B = B * 10$	Read(F)
Write(B)	Unlock(F)
Lock_X(F)	Lock_S(B)
$F = B * 100$	
Write(F)	
Unlock(F)	
Unlock(B)	
	Read(B)
	Print(F+B)
	Unlock(B)

- Which locks have to wait?
- Is this conflict serializable?




Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	Lock_S(F)
B = B*10	Read(F)
Write(B)	Unlock(F)
Lock_X(F)	Lock_S(B)
F = B*100	
Write(F)	
Unlock(F)	
Unlock(B)	
	Read(B)
	Print(F+B)
	Unlock(B)

- 2PL?
- Strict 2PL?
- No deadlock?

Worksheet - 3

T1	T2
Lock_X(B)	
Read(B)	Lock_S(F)
$B = B * 10$	Read(F)
Write(B)	Unlock(F)
Lock_X(F)	Lock_S(B)
$F = B * 100$	
Write(F)	
Unlock(F)	
Unlock(B)	
	Read(B)
	Print(F+B)
	Unlock(B)

- 2PL? 
- Strict 2PL? 
- No deadlock? 

- There were some slides here that were incorrect. See “Discussion 8 Corrections” for more.

Deadlock Detection

- “Waits-for” Graph:
 - edge from T_i to T_j if i is waiting for j
- Deadlock if cycle in graph

Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

- What needs to wait?

T1

T2

T3

T4

Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

- What needs to wait?

T1 → T2

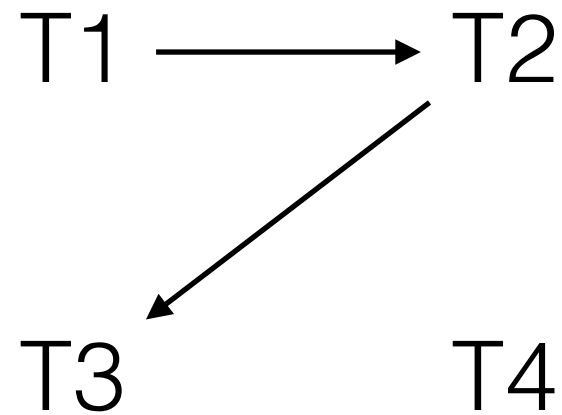
T3

T4

Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

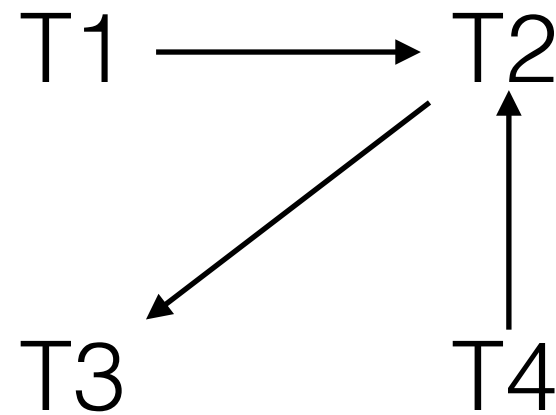
- What needs to wait?



Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

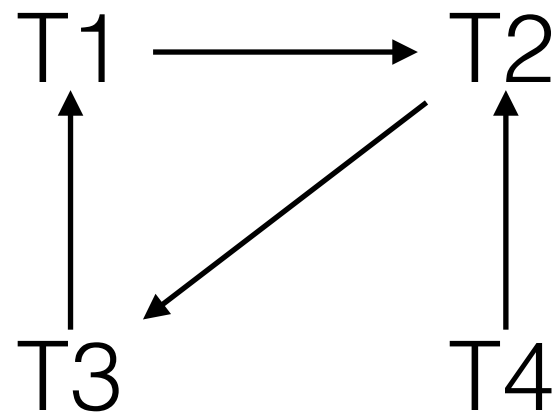
- What needs to wait?



Deadlock Detection

T1	S(A)	S(D)		S(B)					
T2			X(B)				X(C)		
T3					S(D)	S(C)			X(A)
T4								X(B)	

- What needs to wait?



Deadlock Avoidance

- Priorities based on timestamp
- Wait-Die: if T_i higher, wait for T_j , else *abort T_i*
- Wound-Wait: if T_i higher, *T_j aborts*, else wait for T_j